What the Table Shows

PoisonedRAG, Table 4 (NQ)

PoisonedRAG (Black-Box)	**0.97**
Prompt Injection Attack	0.62
Disinformation Attack	0.69
Naive Attack	0.03
Corpus Poisoning Attack	**0.01**

The sentence a reader forms

“Corpus poisoning does not work.”

Both numbers are correct. Neither supports that sentence.

What Each Number Counts

Shared by both

Natural Questions corpus • pre-trained Contriever retriever

PoisonedRAG: 0.01

counts a target question answered with the attacker's answer
over 5 retrieved texts

unit fraction of targets

Corpus Poisoning: 84.2%

counts an adversarial passage entering the retrieved set
over top-20 retrieved

unit percentage of queries

Corpus poisoning never chooses an answer, so it cannot score on an answer-manipulation metric.

What the Audit Returns

Coordinates, both sides

C1	evidence state	match
C2	origin / capability	match
C4	outcome / unit	**mismatch**
C5	policy / utility	match

Verdict

undefined

no operating point makes 1% and 84.2% answer one question

What may be reported

These two attacks cannot be ranked by these two numbers — and the benchmarked paper said so.